

Link Github: https://github.com/paulbboone/DataMining_ThucHanh


HOMEWORK LAB 04

```
[14]: #Import necessary modules
      from sklearn import datasets
      import matplotlib.pyplot as plt

      #Load the digits dataset: digits
      digits = datasets.load_digits()
```

```
[15]: #Display image 1010
      plt.imshow(digits.images[1010], cmap=plt.cm.gray_r, interpolation='nearest')
```

t[15]: <matplotlib.image.AxesImage at 0x17037db0880>



In []:

```
In [16]: from sklearn.model_selection import train_test_split
      X = digits.data
      y = digits.target

      #Split into training and test set
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42, stratify=y)
```

In []:

```
In [17]: #Import necessary modules
      from sklearn.neighbors import KNeighborsClassifier

      import numpy as np

      #Create a k-NN classifier with 3 neighbors: knn
      knn = KNeighborsClassifier(n_neighbors=3)

      #Fit the classifier to the training data
      knn.fit(X_train, y_train)

      #Print the accuracy
      print("Accuracy: {}".format(knn.score(X_test, y_test)))

      Accuracy: 0.9861111111111112
```

```

18]: #Setup arrays to store train and test accuracies
neighbors = np.arange(1, 9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))

#Loop over different values of k

for i, k in enumerate(neighbors):
    #Setup a k-NN Classifier with k neighbors: knn
    knn = KNeighborsClassifier(n_neighbors=k)


    #Fit the classifier to the training data
    knn.fit(X_train, y_train)

    #Compute accuracy on the testing set
    train_accuracy[i] = knn.score(X_train, y_train)

    #Compute accuracy on the training set
    test_accuracy[i] = knn.score(X_test, y_test)

#Generate plot
plt.title('k-NN: Varying Number of Neighbors')
plt.plot(neighbors, test_accuracy, label = 'Testing Accuracy')
plt.plot(neighbors, train_accuracy, label = 'Training Accuracy')
plt.legend()
plt.xlabel('Number of Neighbors')
plt.ylabel('Accuracy')
plt.show()

```



In [21]: `!pip install torch`

```
Requirement already satisfied: torch in c:\anaconda3\lib\site-packages (2.0.0)
Requirement already satisfied: sympy in c:\anaconda3\lib\site-packages (from torch) (1.10.1)
Requirement already satisfied: filelock in c:\anaconda3\lib\site-packages (from torch) (3.6.0)
Requirement already satisfied: networkx in c:\anaconda3\lib\site-packages (from torch) (2.8.4)
Requirement already satisfied: typing-extensions in c:\anaconda3\lib\site-packages (from torch) (4.3.0)
Requirement already satisfied: jinja2 in c:\anaconda3\lib\site-packages (from torch) (2.11.3)
Requirement already satisfied: MarkupSafe in c:\anaconda3\lib\site-packages (from jinja2->torch) (2.0.1)
Requirement already satisfied: mpmath<=0.19 in c:\anaconda3\lib\site-packages (from sympy->torch) (1.2.1)
```

In [22]: `from __future__ import print_function
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.autograd import variable`

In [24]: `!pip install torchvision`

```
Collecting torchvision
  Downloading torchvision-0.15.1-cp39-cp39-win_amd64.whl (1.2 MB)
    ----- 1.2/1.2 MB 1.5 MB/s eta 0:00:00
Requirement already satisfied: torch==2.0.0 in c:\anaconda3\lib\site-packages (from torchvision) (2.0.0)
Requirement already satisfied: requests in c:\anaconda3\lib\site-packages (from torchvision) (2.28.1)
Requirement already satisfied: numpy in c:\anaconda3\lib\site-packages (from torchvision) (1.21.5)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in c:\anaconda3\lib\site-packages (from torchvision) (9.2.0)
Requirement already satisfied: filelock in c:\anaconda3\lib\site-packages (from torch==2.0.0->torchvision) (3.6.0)
Requirement already satisfied: typing-extensions in c:\anaconda3\lib\site-packages (from torch==2.0.0
```

In [25]: `from torchvision import datasets, transforms
mnist = datasets.MNIST(root='.', train=True, download=True)`

```
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to .\MNIST\raw\train-images-idx3-ubyte.gz
```

```
100%|#####| 9912422/9912422 [00:06<00:00, 1493800.13it/s]
```

```
Extracting .\MNIST\raw\train-images-idx3-ubyte.gz to .\MNIST\raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to .\MNIST\raw\train-labels-idx1-ubyte.gz
```

```
100%|#####| 28881/28881 [00:00<00:00, 28903768.51it/s]
```

```
Extracting .\MNIST\raw\train-labels-idx1-ubyte.gz to .\MNIST\raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to .\MNIST\raw\t10k-images-idx3-ubyte.gz
```

```
100%|#####| 1648877/1648877 [00:00<00:00, 1836942.84it/s]
```

```
Extracting .\MNIST\raw\t10k-images-idx3-ubyte.gz to .\MNIST\raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to .\MNIST\raw\t10k-labels-idx1-ubyte.gz
```

```
100%|#####| 4542/4542 [00:00<?, ?it/s]
```

```
Extracting .\MNIST\raw\t10k-labels-idx1-ubyte.gz to .\MNIST\raw
```

```
In [26]: print("Number of training example", mnist.train_data.shape)
print("Image information", mnist[0])
```


Number of training example torch.Size([60000, 28, 28])

Image information (<PIL.Image.Image image mode=L size=28x28 at 0x17037AF8C40>, 5)

C:\anaconda3\lib\site-packages\torchvision\datasets\mnist.py:75: UserWarning: train_data has been renamed data
warnings.warn("train_data has been renamed data")

```
In [27]: import matplotlib.pyplot as plt
%matplotlib inline
plt.imshow(mnist[0][0])
```

Out[27]: <matplotlib.image.AxesImage at 0x1703b139340>




```
In [38]: class Net(nn.Module):
        def __init__(self):
            super(Net, self).__init__()
            self.fully = nn.Sequential(
                nn.Linear(28*28, 10)
            )
        def forward(self, x):
            x = x.view([-1,28*28])
            x = self.fully(x)
            x = F.log_softmax(x, dim=1)
            return x
```

```
In [39]: train_loader = torch.utils.data.DataLoader(datasets.MNIST(root='.', train=True, transform=transforms.Compose([
test_loader = torch.utils.data.DataLoader(datasets.MNIST(root='.', train=False, transform=transforms.Compose([
```

```
In [40]: def train():
        learning_rate = 1e-3
        num_epochs = 3

        net = Net()
        optimizer = torch.optim.Adam(net.parameters(), lr=learning_rate)

        for epoch in range(num_epochs):
            for batch_idx, (data, target) in enumerate(train_loader):
                output = net(data)

                loss = F.nll_loss(output, target)
                optimizer.zero_grad()
                loss.backward()
                optimizer.step()

                if batch_idx % 100 == 0:
                    print('Epoch = %f. Batch = %s. Loss = %s' % (epoch, batch_idx, loss.item()))

        return net
```

```
In [41]: net = train()
```

```
Epoch = 0.000000. Batch = 0. Loss = 2.359036684036255
Epoch = 0.000000. Batch = 100. Loss = 0.8158131241798401
Epoch = 0.000000. Batch = 200. Loss = 0.6744622588157654
Epoch = 0.000000. Batch = 300. Loss = 0.6809446215629578
Epoch = 0.000000. Batch = 400. Loss = 0.4705858826637268
Epoch = 0.000000. Batch = 500. Loss = 0.551677405834198
Epoch = 0.000000. Batch = 600. Loss = 0.23497943580150604
Epoch = 0.000000. Batch = 700. Loss = 0.37437042593955994
Epoch = 0.000000. Batch = 800. Loss = 0.32086634635925293
Epoch = 0.000000. Batch = 900. Loss = 0.46935322880744934
Epoch = 1.000000. Batch = 0. Loss = 0.2193935215473175
Epoch = 1.000000. Batch = 100. Loss = 0.5654202103614807
Epoch = 1.000000. Batch = 200. Loss = 0.4545310139656067
Epoch = 1.000000. Batch = 300. Loss = 0.24608594179153442
Epoch = 1.000000. Batch = 400. Loss = 0.37739771604537964
Epoch = 1.000000. Batch = 500. Loss = 0.38584208488464355
Epoch = 1.000000. Batch = 600. Loss = 0.21674884855747223
Epoch = 1.000000. Batch = 700. Loss = 0.2655832767486572
Epoch = 1.000000. Batch = 800. Loss = 0.3925359547138214
Epoch = 1.000000. Batch = 900. Loss = 0.2943882942199707
Epoch = 2.000000. Batch = 0. Loss = 0.2143695056438446
Epoch = 2.000000. Batch = 100. Loss = 0.3782268464565277
Epoch = 2.000000. Batch = 200. Loss = 0.36743879318237305
Epoch = 2.000000. Batch = 300. Loss = 0.39899033308029175
Epoch = 2.000000. Batch = 400. Loss = 0.25580480694770813
Epoch = 2.000000. Batch = 500. Loss = 0.35438457131385803
Epoch = 2.000000. Batch = 600. Loss = 0.22726872563362122
Epoch = 2.000000. Batch = 700. Loss = 0.3351910710334778
Epoch = 2.000000. Batch = 800. Loss = 0.1402580738067627
Epoch = 2.000000. Batch = 900. Loss = 0.2575857937335968
```



```
In [46]: from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split

x_fertility = df['fertility'].values.reshape(-1,1)
y_life = df['life'].values.reshape(-1,1)
prediction_space = np.linspace(min(x_fertility), max(x_fertility)).reshape(-1,1)

x_train, x_test, y_train, y_test = train_test_split(x_fertility, y_life, test_size=0.3, random_state=42)


reg=LinearRegression()

reg.fit(x_train, y_train)
y_predict=reg.predict(prediction_space)

print(reg.score(x_fertility, y_life))

plt.scatter(x_fertility, y_life, color='blue')
plt.plot(prediction_space, y_predict, color='black', linewidth=3)
plt.ylabel('Life Expectancy')
plt.xlabel('Fertility')
plt.show()
```

0.6162438752151917



```
In [47]: features = pd.read_csv('gapminder.csv')
df=pd.read_csv('gapminder.csv')
del features['life']
del features['Region']

y_life = df['life'].values.reshape(-1,1)


x_train, x_test, y_train, y_test = train_test_split(features, y_life, test_size=0.3, random_state=42)

reg_all = LinearRegression()
reg_all.fit(x_train, y_train)

print(reg_all.score(features, y_life))
```

0.8914651485793137

```
In [49]: N=10
m=.9
c=1
x=np.linspace(0, 2*np.pi,N)
y= m*x + c+ np.random.normal(0, .3, x.shape)
plt.figure()
plt.plot(x,y,'o')
plt.xlabel('x')
plt.ylabel('y')
plt.title('2D data (#data =%d)' %N)
plt.show()
```



```
In [52]: from torch.utils.data import Dataset
class MyDataset(Dataset):
    def __init__(self, x, y):
        self.x = x
        self.y = y

    def __len__(self):
        return len(self.x)

    def __getitem__(self, idx):
        sample = {
            'feature': torch.tensor([1, self.x[idx]]),
            'label': torch.tensor([self.y[idx]])
        }
        return sample
```



```
In [53]: dataset=MyDataset(x,y)
         for i in range(len(dataset)):
             sample=dataset[i]
             print(i, sample['feature'], sample['label'])

0 tensor([1., 0.], dtype=torch.float64) tensor([0.9466], dtype=torch.float64)
1 tensor([1.0000, 0.6981], dtype=torch.float64) tensor([2.2234], dtype=torch.float64)
2 tensor([1.0000, 1.3963], dtype=torch.float64) tensor([2.2778], dtype=torch.float64)
3 tensor([1.0000, 2.0944], dtype=torch.float64) tensor([2.5913], dtype=torch.float64)
4 tensor([1.0000, 2.7925], dtype=torch.float64) tensor([3.4802], dtype=torch.float64)
5 tensor([1.0000, 3.4907], dtype=torch.float64) tensor([4.2653], dtype=torch.float64)
6 tensor([1.0000, 4.1888], dtype=torch.float64) tensor([4.5442], dtype=torch.float64)
7 tensor([1.0000, 4.8869], dtype=torch.float64) tensor([5.1212], dtype=torch.float64)
8 tensor([1.0000, 5.5851], dtype=torch.float64) tensor([6.3011], dtype=torch.float64)
9 tensor([1.0000, 6.2832], dtype=torch.float64) tensor([7.2109], dtype=torch.float64)
```

```
In [56]: from torch.utils.data import DataLoader

dataset=MyDataset(x,y)
batch_size=4
shuffle=True
num_workers=0
dataloader= DataLoader(dataset, batch_size=batch_size, shuffle=shuffle, num_workers=num_workers)
```

```
In [57]: import pprint as pp
         for i_batch, samples in enumerate(dataloader):
             print('\nbatch# = %s' % i_batch)
             print('samples:')
             pp.pprint(samples)
```

```
batch# = 0
samples:
{'feature': tensor([[1.0000, 2.0944],
                    [1.0000, 6.2832],
                    [1.0000, 5.5851],
                    [1.0000, 0.6981]], dtype=torch.float64),
 'label': tensor([[2.5913],
                  [7.2109],
                  [6.3011],
                  [2.2234]], dtype=torch.float64)}
```

```
batch# = 1
samples:
{'feature': tensor([[1.0000, 4.1888],
                    [1.0000, 4.8869],
                    [1.0000, 0.0000],
                    [1.0000, 3.4907]], dtype=torch.float64),
 'label': tensor([[4.5442],
                  [5.1212],
                  [0.9466],
                  [4.2653]], dtype=torch.float64)}
```

```
batch# = 2
samples:
{'feature': tensor([[1.0000, 1.3963],
                    [1.0000, 2.7925]], dtype=torch.float64),
 'label': tensor([[2.2778],
                  [3.4802]], dtype=torch.float64)}
```

```
In [59]: import torch.nn as nn
import torch.nn.functional as F
class MyModel(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(MyModel, self).__init__()
        self.linear = nn.Linear(input_dim, output_dim)

    def forward(self, x):
        out=self.linear(x)
        return out
```

```
In [60]: input_dim=2
output_dim=1

model = MyModel(input_dim, output_dim)
```

```
In [61]: cost=nn.MSELoss()
```

```
In [90]: num_epochs=10
         lr_rate=0.01
         optimiser= torch.optim.SGD(model.parameters(), lr=lr_rate)

         dataset= MyDataset(x,y)
         batch_size=4
         shuffle=True
         num_workers=0
         training_sample_generator= DataLoader(dataset, batch_size=batch_size, shuffle=shuffle, num_workers=num_

         for epoch in range(num_epochs):
             print('Epoch =%s' % epoch)
             for batch_i, samples in enumerate(training_sample_generator):
                 predictions = model(samples['feature'])
                 error = cost(predictions, samples['label'])
                 print('\tBatch = %s, Error = %s' % (batch_i, error.item()))

                 optimiser.zero_grad()

                 error.backward()

                 optimiser.step()

             print('\tBatch = %s, Error = %s' % (batch_i, error.item()))
```

```
C:\anaconda3\lib\site-packages\torch\nn\modules\module.py in _call_impl(self, *args, **kwargs)
    1499         or _global_backward_pre_hooks or _global_backward_hooks
    1500         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1501         return forward_call(*args, **kwargs)
    1502         # Do not call functions when jit is used
    1503         full_backward_hooks, non_full_backward_hooks = [], []

~\AppData\Local\Temp\ipykernel_22480\1038057100.py in forward(self, x)
      7
      8     def forward(self, x):
----> 9         out=self.linear(x)
      10         return out

C:\anaconda3\lib\site-packages\torch\nn\modules\module.py in _call_impl(self, *args, **kwargs)
    1499         or _global_backward_pre_hooks or _global_backward_hooks
    1500         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1501         return forward_call(*args, **kwargs)
    1502         # Do not call functions when jit is used
    1503         full_backward_hooks, non_full_backward_hooks = [], []

C:\anaconda3\lib\site-packages\torch\nn\modules\linear.py in forward(self, input)
    112
    113     def forward(self, input: Tensor) -> Tensor:
--> 114         return F.linear(input, self.weight, self.bias)
    115
    116     def extra_repr(self) -> str:


RuntimeError: mat1 and mat2 must have the same dtype
```

```
In [65]: x_for_plotting = np.linspace(0, 2*np.pi, 1000)
         design_matrix = torch.tensor(np.vstack([np.ones(x_for_plotting.shape), x_for_plotting])).T, dtype=torch.
         print('Design matrix shape:', design_matrix.shape)

         y_for_plotting = model.forward(design_matrix)
         print('y_for_plotting shape:', y_for_plotting.shape)

         Design matrix shape: torch.Size([1000, 2])
         y_for_plotting shape: torch.Size([1000, 1])
```

```
In [66]: plt.figure()
plt.plot(x,y,'o')
plt.plot(x_for_plotting, y_for_plotting.data.numpy(), 'r-')
plt.xlabel('x')
plt.ylabel('y')
plt.title('2D data (#data = %d)' % N)
plt.show()
```



In [68]: # RECOMMENDATION SYSTEMS

```
In [73]: def user_cf(M, metric='cosine'):
    pred = np.copy(M)
    n_users, n_items = M.shape
    avg_ratings = np.nanmean(M, axis=1)
    sim_users = sim_matrix(M, 'user', metric)
    for i in range(n_users):
        for j in range(n_items):
            if np.isnan(M[i,j]):
                pred[i,j] = avg_ratings[i] + np.nansum(sim_users[i] * (M[:,j] - avg_ratings)) / sum(sim_users[i])
    return pred
```

```
In [74]: def item_cf(M, metric='cosine'):
    pred = np.copy(M)
    n_users, n_items = M.shape
    avg_ratings = np.nanmean(M, axis=0)
    sim_items = sim_matrix(M, 'item', metric)
    for i in range(n_users):
        for j in range(n_items):
            if np.isnan(M[i,j]):
                pred[i,j] = avg_ratings[j] + np.nansum(sim_items[j] * (M[i,:] - avg_ratings)) / sum(sim_items[j])
    return pred
```

```
In [82]: def evaluateRS(M, M_result, algorithm, similarity):
    evaluateRS(M, M_result, 'user_cf', 'cosine')
    evaluateRS(M, M_result, 'user_cf', 'correlation')
    evaluateRS(M, M_result, 'item_cf', 'cosine')
    evaluateRS(M, M_result, 'item_cf', 'correlation')
```

```
In [79]: results=[]
for method in ['user_cf', 'item_cf']:
    for metric in ['cosine', 'correlation']:
        rank_acc = evaluate_rank(M, M_result, method, metric)
        results += ["Rank accuracy of {0} with {1} metric: {2}".format(method[1], metric, rank_acc)]
print("\n".join(results))
```

```
-----
NameError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_22480\3831842885.py in <module>
      2 for method in ['user_cf', 'item_cf']:
      3     for metric in ['cosine', 'correlation']:
----> 4         rank_acc = evaluate_rank(M, M_result, method, metric)
      5         results += ["Rank accuracy of {0} with {1} metric: {2}".format(method[1], metric, rank_acc)]
      6 print("\n".join(results))

NameError: name 'evaluate_rank' is not defined
```

```
In [81]: def evaluate_rank(M, M_result, method, metric):
    results=[]
    for method in ['user_cf', 'item_cf']:
        for metric in ['cosine', 'correlation']:
            rank_acc = evaluate_rank(M, M_result, method, metric)
            results += ["Rank accuracy of {0} with {1} metric: {2}".format(method[1], metric, rank_acc)]
    print("\n".join(results))
```